

Multi-Cloud Decision & Onboarding Wizard

User Guide (v1.5.4)

Purpose. This wizard helps you choose a practical cloud landing pattern, migration approach, and environment footprint for a single workload or initiative, across Azure, AWS, GCP or OCI.

What it produces. Based on your answers, it generates a summary, design guidance for compute/data/integration/ops/migration, a sizing & environments view, and a ready-to-paste implementation playbook.

How it runs. It is a single HTML file that runs fully in the browser. No install, no backend, no data is stored.

1. Opening the Wizard

1. Unzip the ZIP file you were given (for example, *cloud-decision-kit-v1.5.2.zip*).
2. Inside the unzipped folder, open *cloud-decision-kit/multi-cloud-decision-matrix-v5.4.html*.
3. Your default browser will open the wizard. Use a modern browser such as Edge, Chrome, Firefox or Safari.

2. Screen Layout

The screen has three key elements:

- **Cloud selector (top-right).** Choose Azure, AWS, GCP, or OCI. All recommendations and service names are tailored to the chosen cloud.
- **Left card – 4-step input wizard.** Step 1 collects initiative basics. Step 2 collects path details and data profile. Step 3 covers non-functional requirements and migration tooling. Step 4 collects sizing and environment details.
- **Right card – Results.** After you press Finish on Step 4, the right side displays the recommended patterns, a sizing & environment footprint, and an implementation playbook that you can copy into Word, Confluence, or slideware.

3. Basic Workflow

For each workload or initiative, follow this standard workflow:

4. Select the target cloud in the top-right (Azure, AWS, GCP, or OCI).
5. Walk through Steps 1–4 from top to bottom, filling in answers as best you can.
Soft validation will gently warn you if something important is missing but will not block you.
6. On Step 4, press Finish to generate the recommendations and playbook.
7. Review the right-hand results, adjust any inputs that feel wrong, and press Finish again to refresh.
8. Copy the Implementation playbook section into a Word document, email, or Confluence page for sharing.

4. Step 1 – Initiative & Basics

Step 1 describes what you are doing at a high level and who is best positioned to build and run it.

Field	What it means	How it influences the output
Initiative type	Is this a new service, a change to an existing service, maintenance work, or a migration?	Determines which path details you will see in Step 2 and whether the guidance focuses on build, change, run, or move.
Workload / initiative name	A human-readable name such as “Retail payments API” or “Claims analytics hub”.	Used in the prose so that you can recognise the workload; does not change the technical pattern.
Architecture type	The dominant runtime architecture: web/API, microservices/containers, batch, event-driven, legacy VMs, or data/analytics platform.	Shapes compute and data guidance (for example, app services vs containers vs batch services vs data platform).
Traffic pattern	Typical and peak traffic shape: low predictable, medium, high 24/7, or	Helps choose between fixed capacity vs autoscaling/serverless

	spiky/bursty.	patterns and feeds into sizing tiers.
Latency sensitivity	How sensitive the workload is to latency: strict, moderate, or relaxed.	Influences network and region guidance and whether “low-latency” architectures are emphasised.
Team strengths	The primary strength of the delivery and operations team: VMs/servers, containers/Kubernetes, PaaS, or serverless.	Leans the design towards platforms the team can realistically operate (for example, more IaaS if the team is VM-heavy).
Short description	A one- or two-sentence description of what the workload does in business terms.	Used to improve the wording in notes and the playbook, not the core technical pattern.

5. Step 2 – Path Details & Data Profile

Step 2 merges two concepts: the delivery path (new build, change, maintenance, or migration) and the data and sector profile. The exact top block changes depending on the Initiative type selected in Step 1.

5.1 Path details (top of Step 2)

The wizard shows one of these four path panels:

- If Initiative type is “New service / greenfield”: choose the kind of service (customer-facing, internal, analytics platform, integration hub, shared platform) and its stage (pilot, new production, scaling). This adjusts emphasis between fast launch vs long-term scalable architecture.
- If Initiative type is “Change to existing service”: choose the dominant change focus (scale/HA, features, compliance, cost, modernization) and optionally describe current pain points. This shifts the recommendations toward the area you most want to improve.
- If Initiative type is “Maintenance / operations”: select the main maintenance theme (patching, performance, incidents, cost optimisation, SLOs and

reporting) and the operational cadence. This heavily shapes the Ops & resilience guidance.

- If Initiative type is “Migration”: choose the migration scope (single app, portfolio, data centre estate, database-only) and the preferred cutover strategy (big bang, phased, blue/green, canary). This drives the migration-focused sections.

5.2 Data & sector (bottom of Step 2)

The data panel is always visible and is one of the most important drivers of recommendations.

- Primary data pattern – relational OLTP, NoSQL, object/file storage, streaming, or analytics/lake/warehouse. This steers database and storage choices.
- Data sensitivity / sector – public, internal, confidential, regulated, or explicit public-sector/private-sector levels. This influences requirements for encryption, private connectivity, region choice, and control wording.
- Write/update pattern – light, balanced, or heavy writes. This helps distinguish transaction-heavy workloads from read-heavy or batch-style designs.
- Geography / users – single region, multi-region, or global. This controls when multi-region or global front-door patterns are recommended.
- Integration pattern – simple HTTP APIs, enterprise messaging, event streaming, or orchestration. This focuses integration guidance on API gateway, message bus, streaming, or workflow engines.
- Compliance / regulatory notes – any specific regimes (for example PCI, GLBA, SOX, FedRAMP). These are woven into the narrative so the playbook explicitly calls out compliance constraints.

6. Step 3 – Non-functional & Migration Tooling

Step 3 locks in non-functional requirements (SLOs, DR, time-to-market, operations maturity) and how you will move the workload from its current environment.

6.1 Reliability & SLOs

- Business criticality – tier 0 (mission critical), tier 1, or tier 2/3. Higher tiers result in stronger HA/DR patterns and more demanding observability.
- Uptime target – approximate target such as ~99.0%, ~99.5%, ~99.9%, or 99.95%+. This aligns with single-region vs multi-AZ vs multi-region recommendations.

- RTO & RPO – how quickly you must recover and how much data loss you can tolerate. These directly drive backup, replication, and DR-region patterns.
- Time-to-market – urgent (weeks), normal (a few months), or longer-term. Tighter timelines often favour managed services and platform choices that reduce toil.
- Ops / SRE maturity – basic, intermediate, or advanced. This influences whether the playbook assumes minimal existing automation or a mature SRE-style practice.

6.2 Source environment & migration approach

- Source environment – on-prem VMware, bare metal, co-location, another cloud, SaaS, or hybrid. This sets expectations around connectivity, landing zones, and migration tooling.
- Migration approach (7R) – rehost, replatform, refactor, repurchase, retain, retire, or relocate. This is the primary driver for the migration and compute recommendations (for example, IaaS-heavy for rehost vs more PaaS/serverless for refactor).

6.3 IaC / automation tooling

Choose the primary infrastructure-as-code or automation approach you intend to use (for example Terraform, Ansible, cloud-native templates, or pipeline-centric tooling). The wizard then frames the implementation playbook around that discipline, suggesting module design, pipelines, and automation practices that match.

7. Step 4 – Sizing & Environments

Step 4 converts your workload description into a simple T-shirt sizing band and an environment footprint matrix. Inputs here are intentionally approximate rather than exact capacity planning numbers.

- Peak concurrent users and peak requests per second – determine a low, moderate, or high traffic band and whether your compute footprint is described as small, medium, or medium-large.
- Total data volume and daily ingest/change volume – drive the data footprint wording and whether heavier ingestion or streaming patterns are suggested.
- Data retention target – steers storage tiering guidance between hot, cool, and archival tiers and highlights the need for lifecycle policies.
- Environments in scope – check the environments you actually plan to build (Dev, Test, Stage, Prod, DR). These become rows in the sizing table.

- Non-prod scale versus prod – describes how large non-prod environments should be relative to production (for example 50% or 25% of prod capacity).
- Regions / sovereign sites – single region, multi-AZ, two regions, or three-plus regions. This changes the network and DR guidance, especially for global or sovereign scenarios.

8. Understanding the Results

After you click Finish on Step 4, the right-hand card populates with structured guidance. Every time you change inputs and click Finish again, the results refresh immediately.

- Summary and pills – a short description line and a set of pills indicating path, data type, tier, cloud, and other key attributes.
- Compute & platform – recommended runtime platform for the chosen cloud (for example, web app service, containers, functions, or IaaS) and when to consider alternatives.
- Data & storage – suggested database, lake or warehouse services, storage tiers, and key patterns for resilience and cost optimisation.
- Integration & messaging – guidance on API gateways, messaging services, event streaming, and orchestration tools appropriate to your selections.
- Ops, resilience & governance – operational practices, monitoring/observability expectations, HA/DR patterns, and governance hooks matched to your tier and maturity.
- Migration & onboarding focus – primary migration approach, supporting tooling, and a high-level sequence of phases for onboarding to the cloud.
- Sizing & environment footprint – traffic and data footprint text plus an environment matrix showing relative size and notes for each environment.
- Implementation playbook – a multi-phase list of actions that you can copy directly into Word as the starting point for your delivery plan.

9. Using the Wizard Across Clouds

You can reuse the same answers to compare clouds. Run the wizard once, then simply change the Cloud selector and click Finish again to see how the recommendations change per provider.

10. Tips & Best Practices

- Do not overthink every field on the first run. Enter reasonable estimates, generate a first draft, then iterate.
- Always record the workload name, date, and cloud beside any exported playbook so you know which run it came from.
- If a recommendation feels too heavy or too light, adjust the relevant input (for example tier, uptime target, or traffic band) and regenerate.
- Save separate copies of the playbook for different clouds or approaches when you are comparing strategies.
- Use the sizing & environment table as a conversation starter for platform and finance teams, not a final capacity plan.